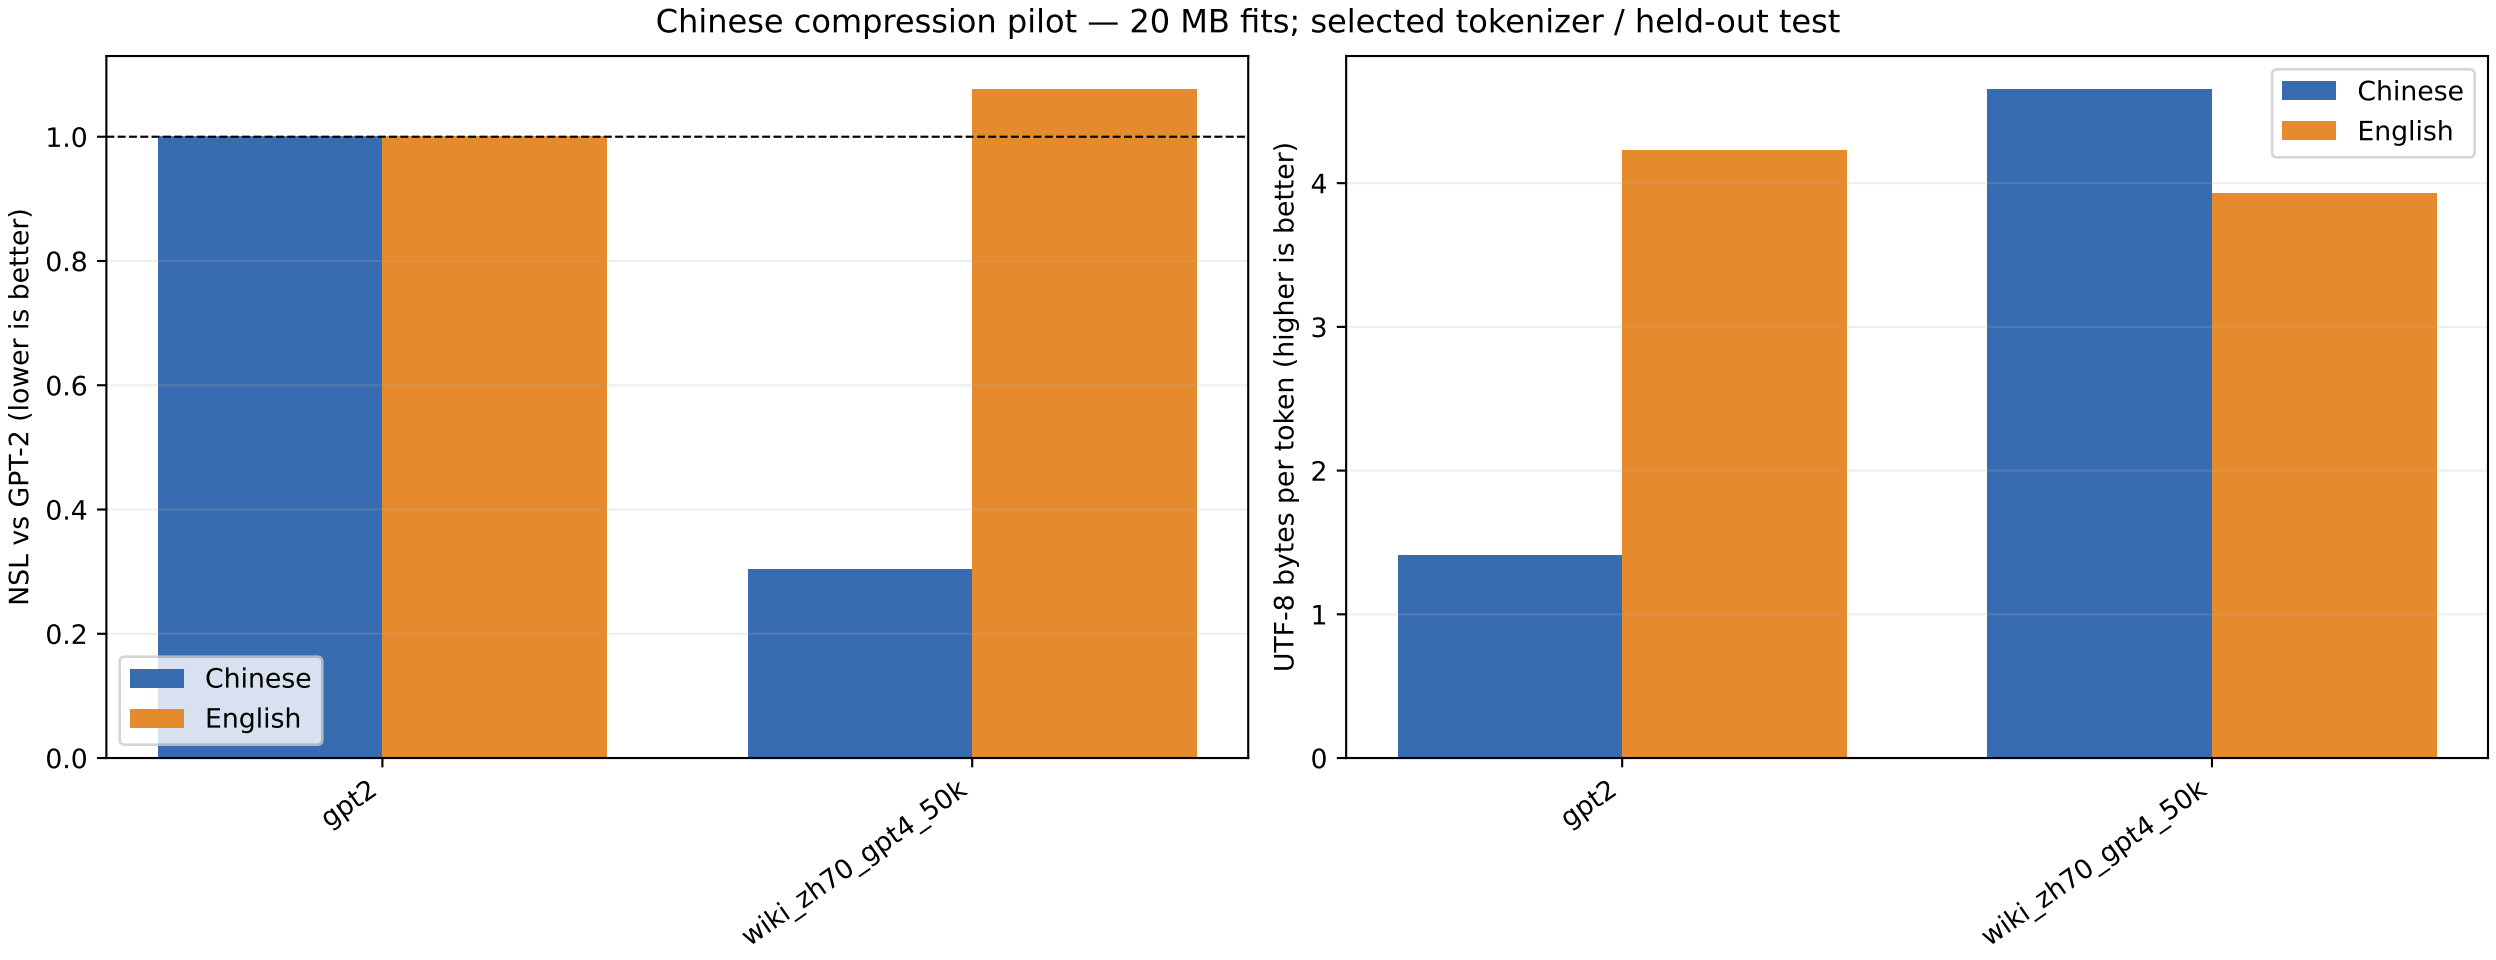
Chinese compression pilot — 20 MB fits; selected tokenizer / held-out test
NSL vs GPT-2 (lower is better)
1.0
0.8
0.6
0.4
0.2
0.0
Chinese
English
gpt2
wiki_zh70_gpt4_50k
UTF-8 bytes per token (higher is better)
4
3
2
1
0
Chinese
English
gpt2
wiki_zh70_gpt4_50k