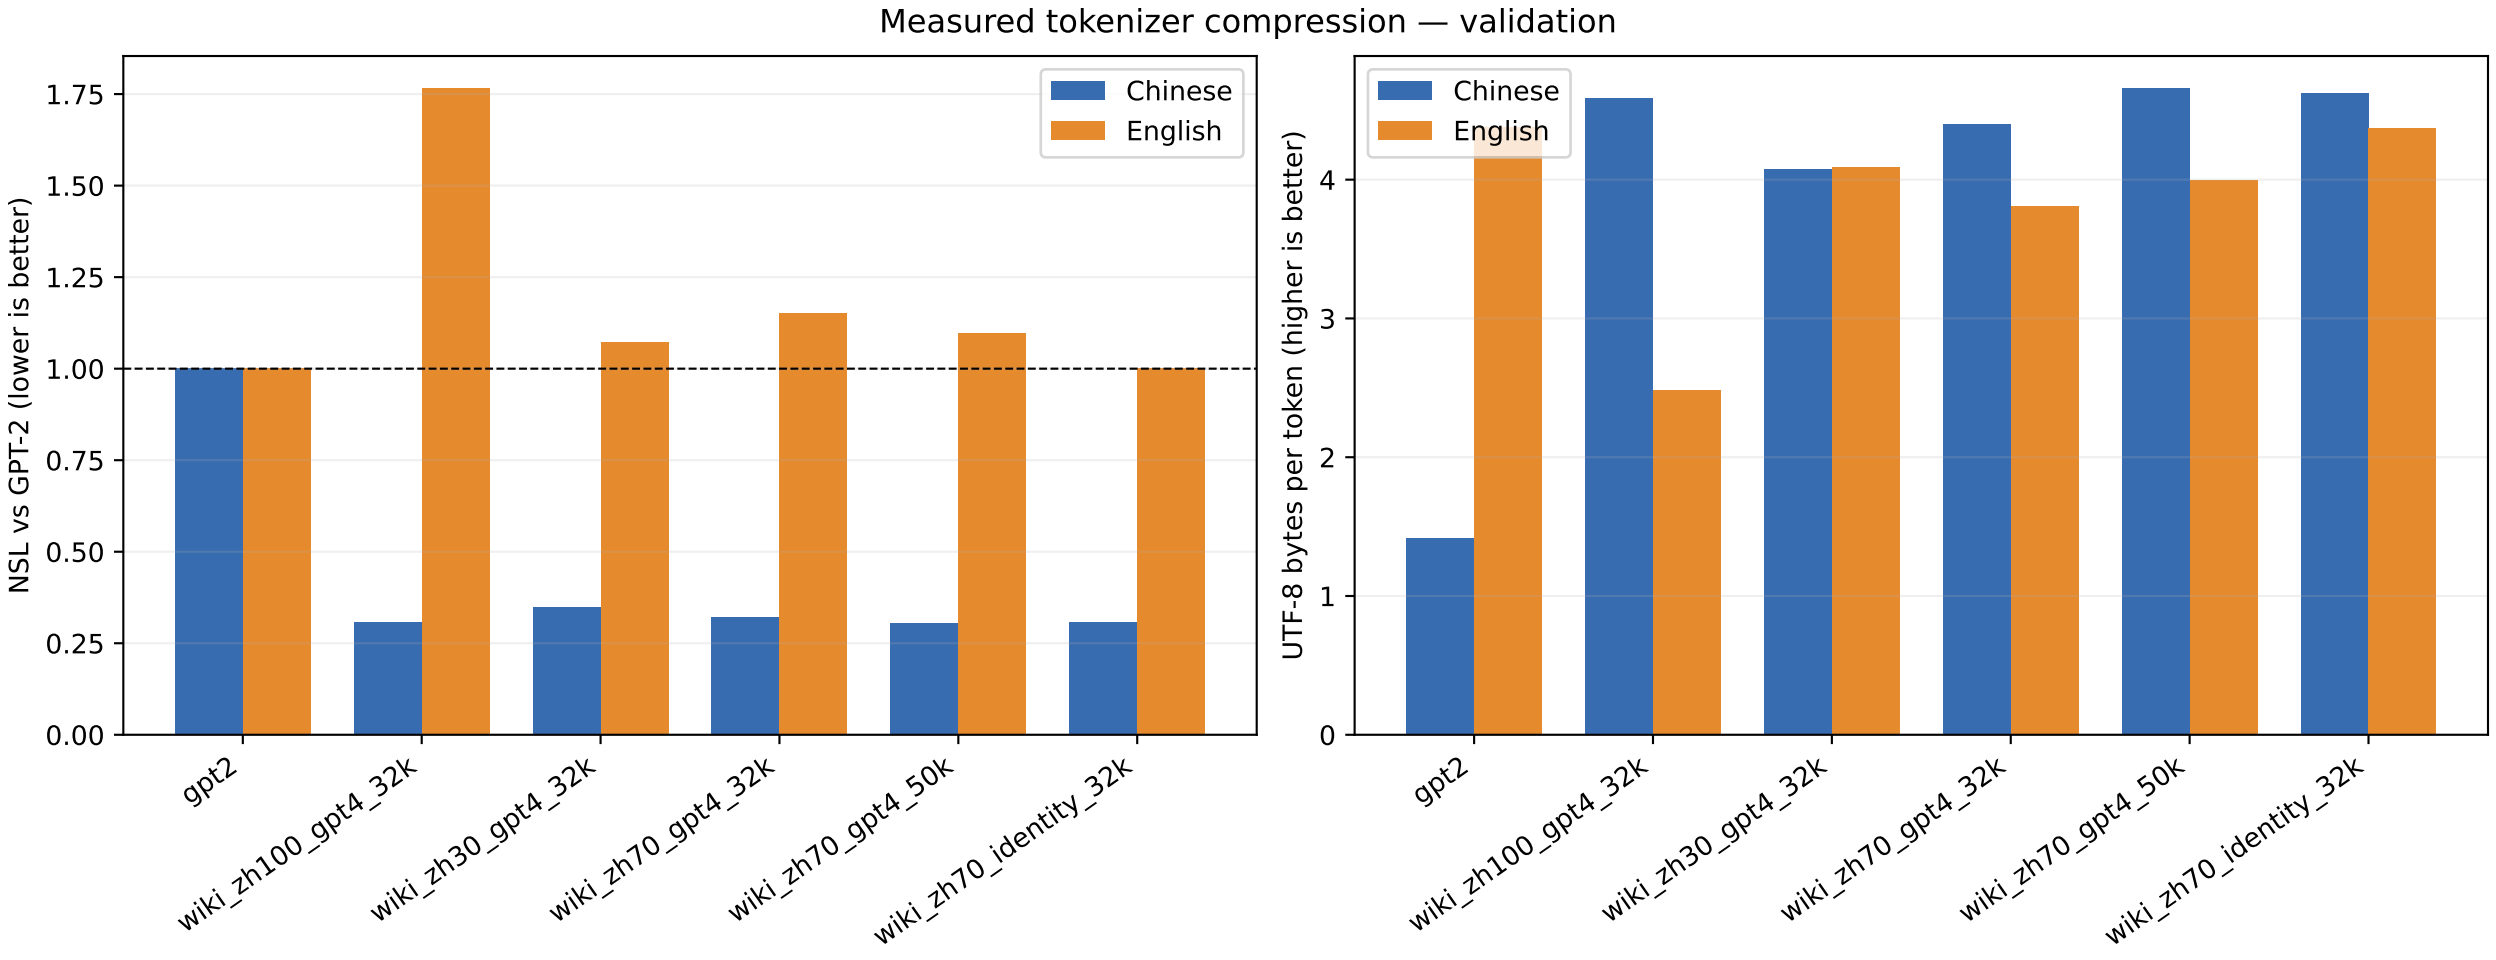
Measured tokenizer compression — validation
NSL vs GPT-2 (lower is better)
UTF-8 bytes per token (higher is better)
Chinese
English
0.00
0.25
0.50
0.75
1.00
1.25
1.50
1.75
0
1
2
3
4
gpt2
wiki_zh100_gpt4_32k
wiki_zh30_gpt4_32k
wiki_zh70_gpt4_32k
wiki_zh70_gpt4_50k
wiki_zh70_identity_32k