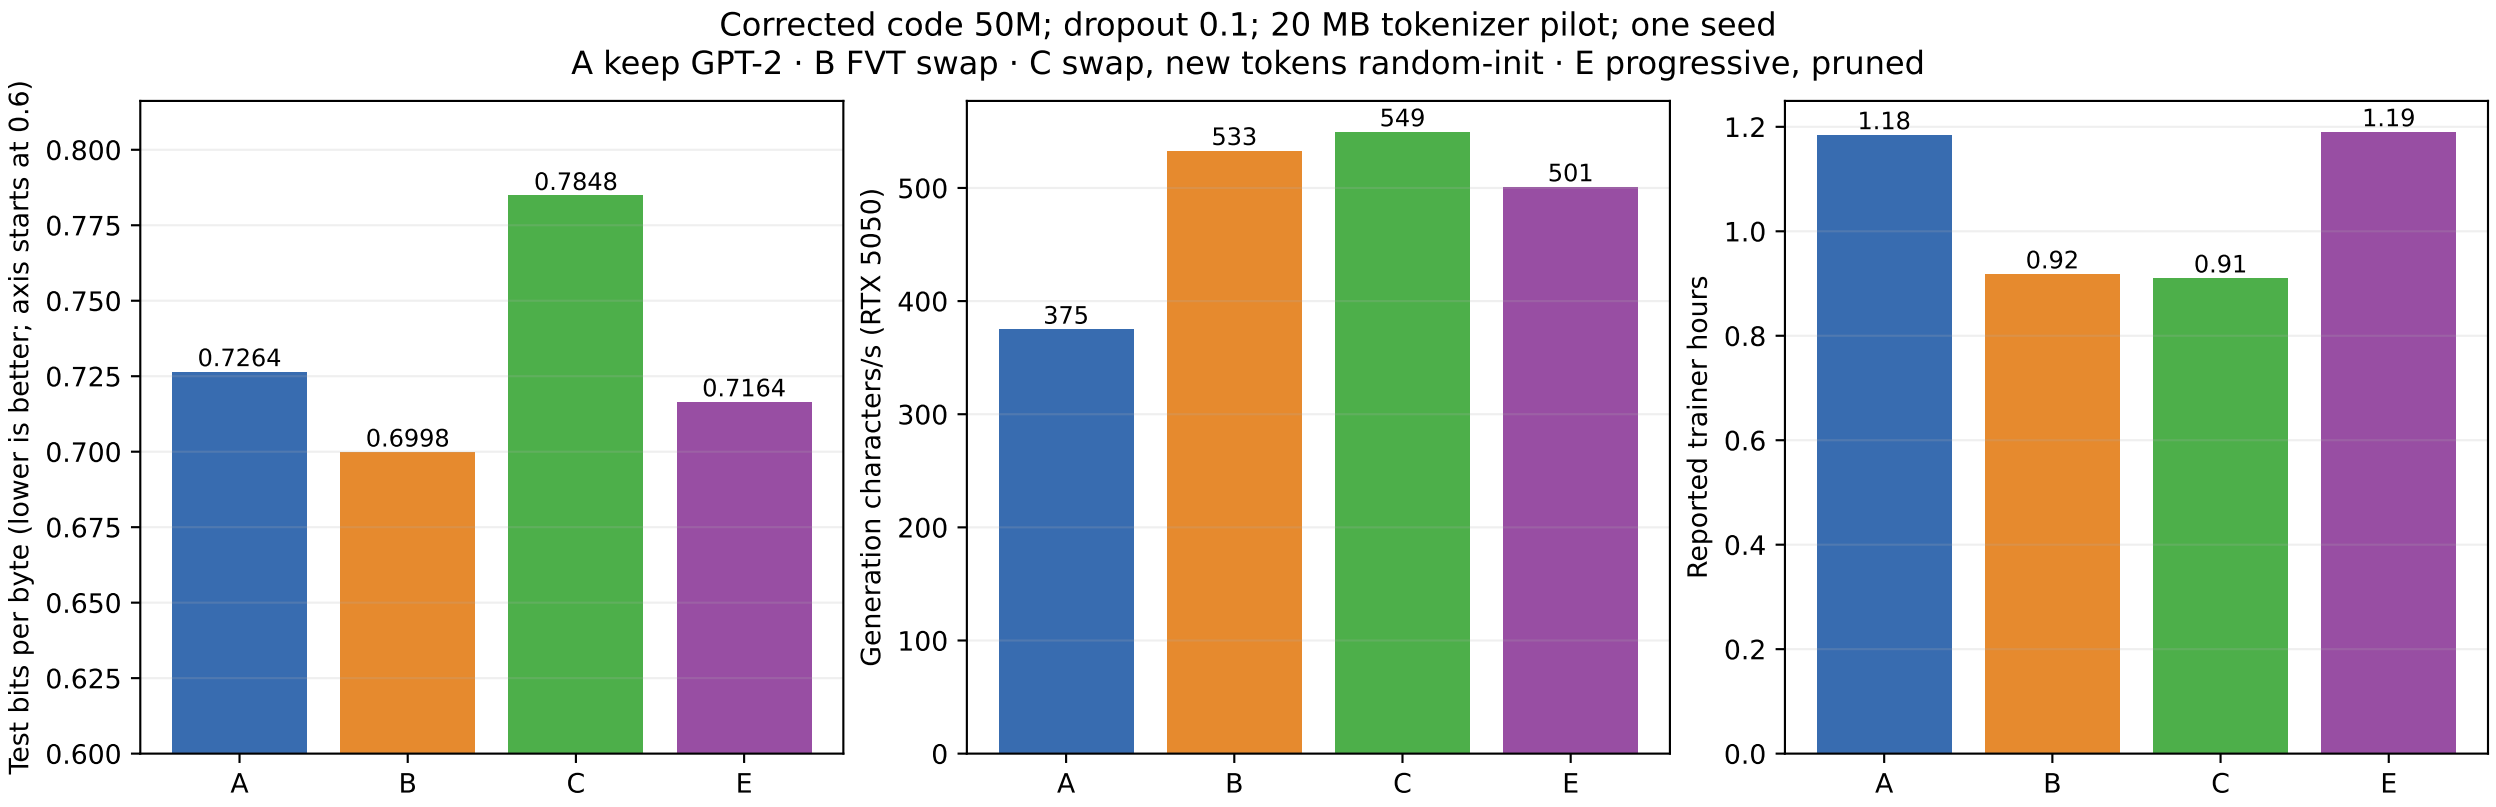
Corrected code 50M; dropout 0.1; 20 MB tokenizer pilot; one seed
A keep GPT-2 · B FVT swap · C swap, new tokens random-init · E progressive, pruned
Test bits per byte (lower is better; axis starts at 0.6)
0.600
0.625
0.650
0.675
0.700
0.725
0.750
0.775
0.800
0.7264
0.6998
0.7848
0.7164
A
B
C
E
Generation characters/s (RTX 5050)
0
100
200
300
400
500
375
533
549
501
A
B
C
E
Reported trainer hours
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.18
0.92
0.91
1.19
A
B
C
E