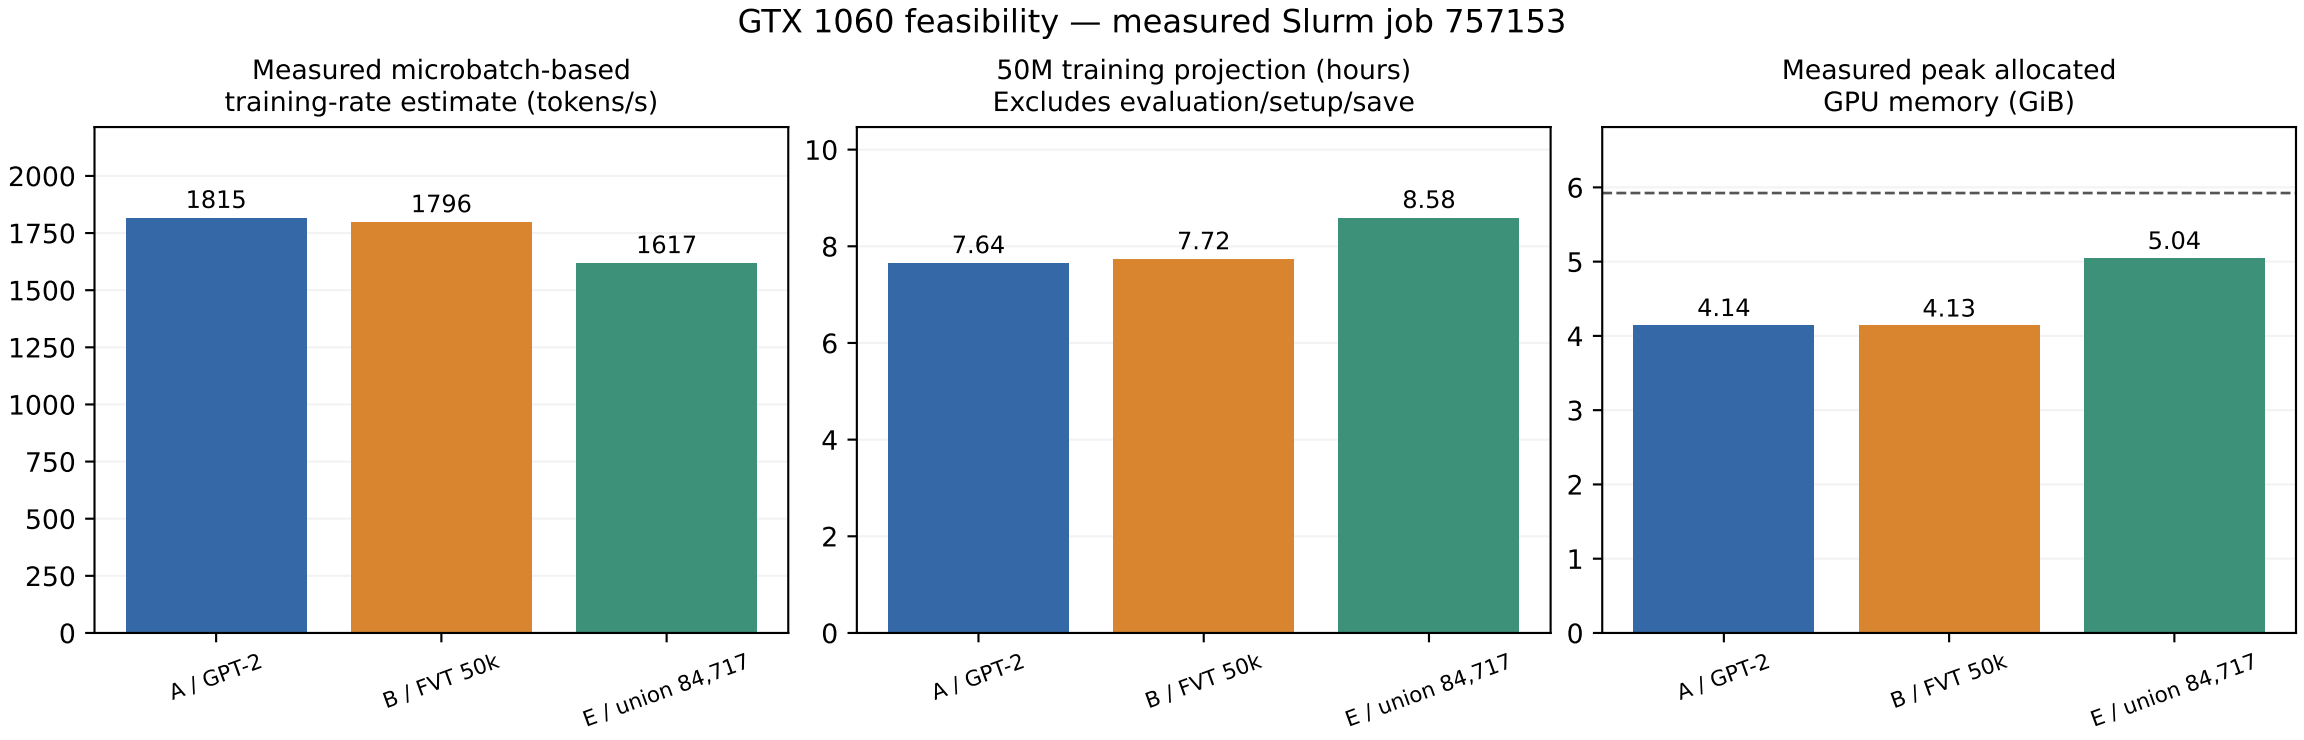
GTX 1060 feasibility — measured Slurm job 757153
Measured microbatch-based
training-rate estimate (tokens/s)
1815
1796
1617
A / GPT-2
B / FVT 50k
E / union 84,717
50M training projection (hours)
Excludes evaluation/setup/save
7.64
7.72
8.58
Measured peak allocated
GPU memory (GiB)
4.14
4.13
5.04
BF16 (no native BF16 hardware), micro-batch 1 × accumulation 64; 65,536 tokens/step.
Four timed microbatches per condition. Diagnostic checkpoints removed. No model-quality claim.