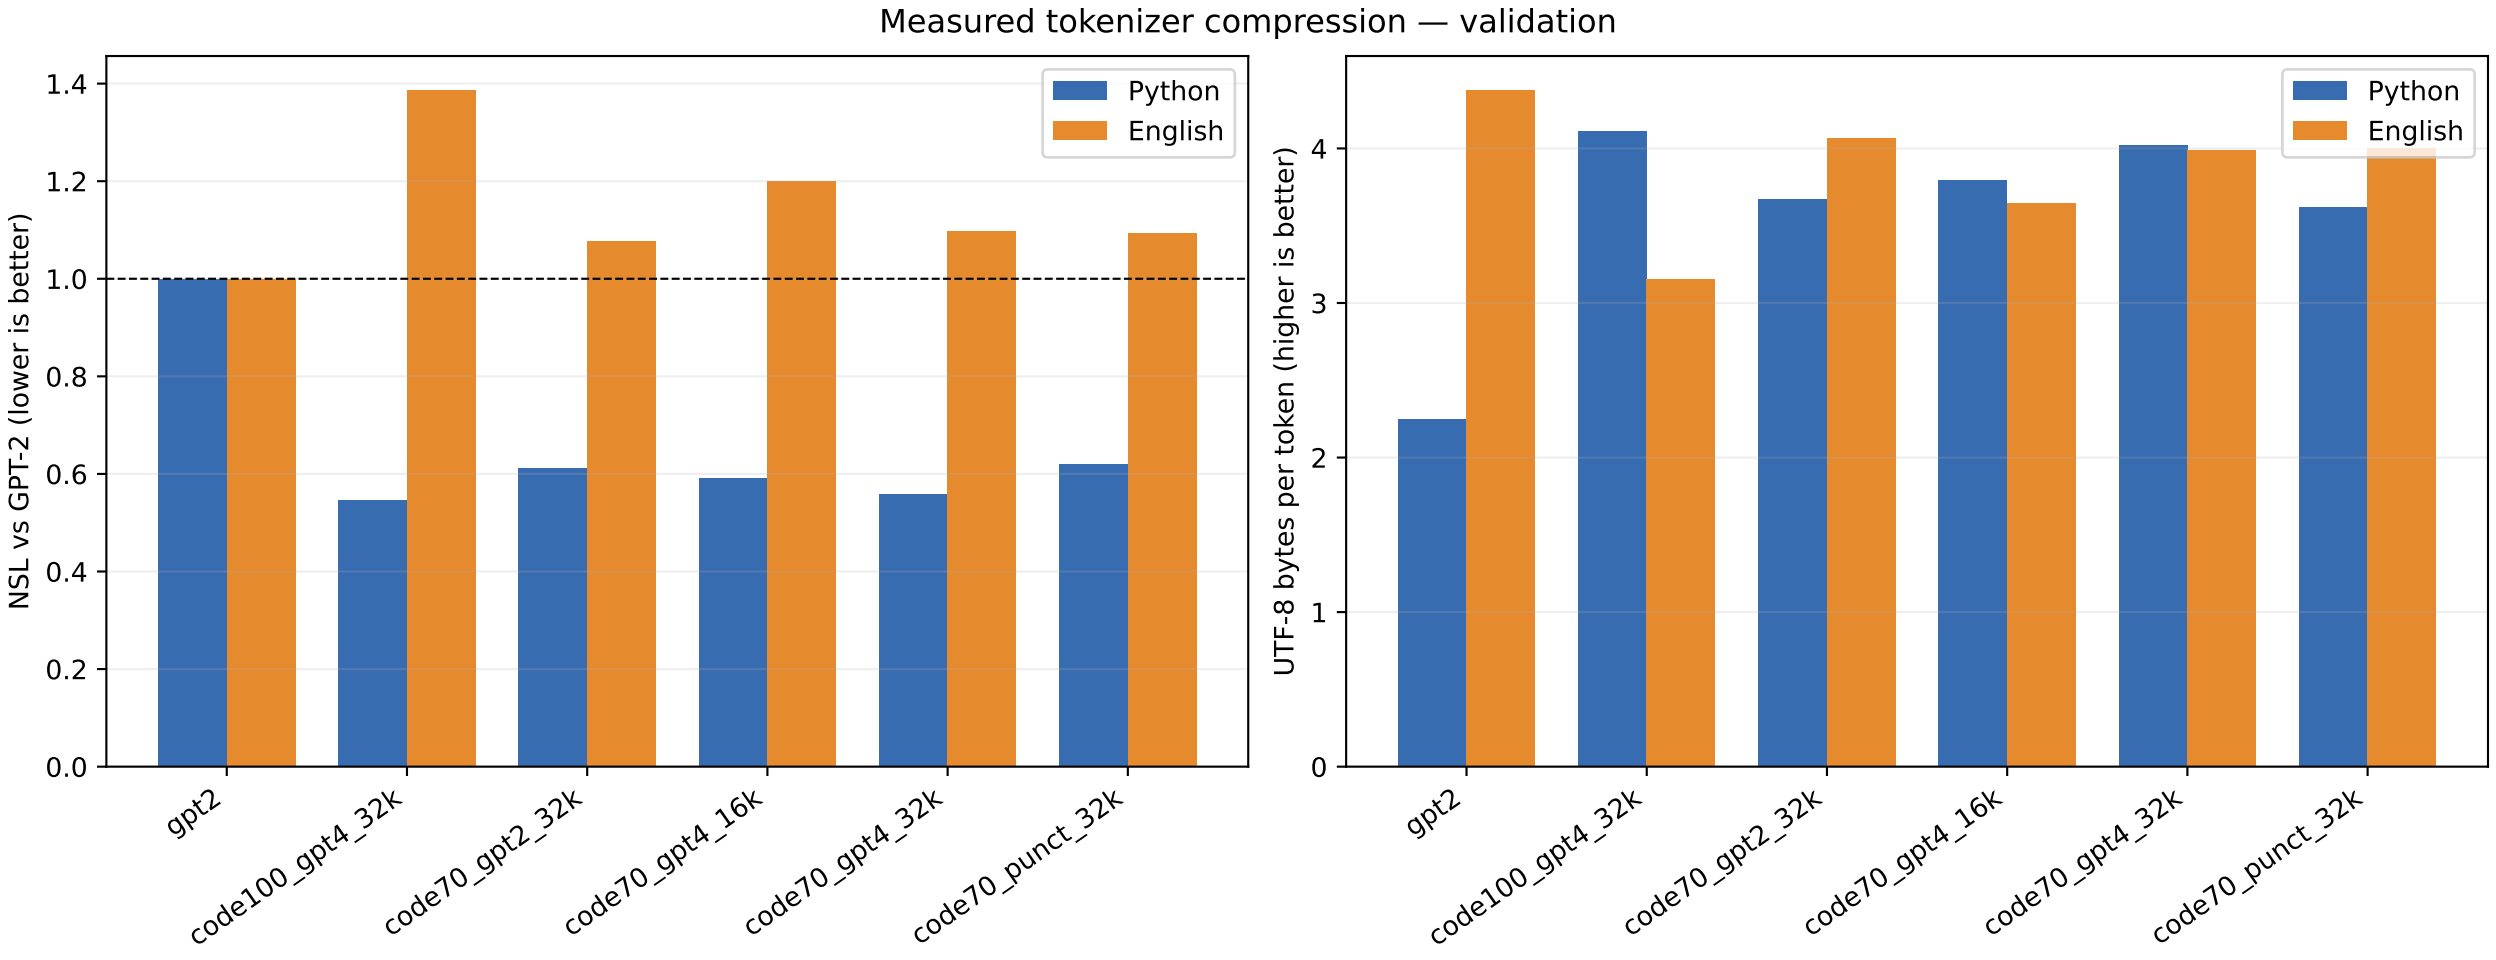
Measured tokenizer compression — validation
NSL vs GPT-2 (lower is better)
UTF-8 bytes per token (higher is better)
Python
English
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
0
1
2
3
4
gpt2
code100_gpt4_32k
code70_gpt2_32k
code70_gpt4_16k
code70_gpt4_32k
code70_punct_32k